

Highlights: Neural Tangent Kernel

Antônio H. Ribeiro, Dave Zachariah, Per Mattsson

Notes about the paper *Neural Tangent Kernel: Convergence and Generalization in Neural Networks* (Jacot et al., 2018).

1 Setup

Let $f(\cdot; \theta) : \mathbb{R}^d \rightarrow \mathbb{R}$ be an almost everywhere differentiable function parametrized by $\theta \in \mathbb{R}^p$. Assume given a training dataset $\{(x_i, y_i)\}_{i=1}^n$ of input and outputs. We define the cost function

$$V(\theta) = \frac{1}{2n} \sum_{i=1}^n (f(x_i; \theta) - y_i)^2. \quad (1)$$

Taking the derivative of V we obtain,

$$\nabla_{\theta} V(\theta) = \frac{1}{n} \sum_{i=1}^n (f(x_i; \theta) - y_i) \nabla_{\theta} f(x_i; \theta). \quad (2)$$

Now assume the parameter θ is estimated using gradient flow, let θ_t be the parameters estimated at the instant t . Then,

$$\frac{d\theta_t}{dt} = -\eta \nabla_{\theta} V(\theta), \quad (3)$$

and the chain rule yields

$$\frac{df(z, \theta_t)}{dt} = \eta \nabla_{\theta} V(\theta)^{\top} \nabla_{\theta} f(z; \theta) = \frac{\eta}{n} \sum_{i=1}^n (f(x_i; \theta) - y_i) (\nabla_{\theta} f(x_i; \theta)^{\top} \nabla_{\theta} f(z; \theta)). \quad (4)$$

We define the Neural Tangent Kernel, $K(\cdot, \cdot; \theta) : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$

$$K(x, z; \theta_t) = \nabla_{\theta} f(x; \theta)^{\top} \nabla_{\theta} f(z; \theta). \quad (5)$$

This is the kernel associated with the feature map $x \mapsto \nabla_{\theta} f(x; \theta)$. It then follows that

$$\frac{df(z, \theta_t)}{dt} = \frac{\eta}{n} \sum_{i=1}^n (f(x_i; \theta) - y_i) K(x_i, z; \theta). \quad (6)$$

2 Model

Here they consider $f(x; \theta) = \tilde{\alpha}^{(\ell)}(x, \theta)$ is the output of a neural network with L layers. that could be defined recursively as:

$$\begin{aligned} \alpha^{(0)}(x; \theta) &= x \\ \tilde{\alpha}^{(\ell+1)}(x; \theta) &= \frac{1}{\sqrt{n_{\ell}}} W^{(\ell)} \alpha^{(\ell)}(x; \theta) + \beta b^{(\ell)} \\ \alpha^{(\ell)}(x; \theta) &= \sigma \left(\tilde{\alpha}^{(\ell)}(x; \theta) \right) \end{aligned} \quad (7)$$

where the nonlinearity σ is applied entrywise and β is a scaling factor. Here $\theta = (W^{(1)}, b^{(1)}, W^{(2)}, b^{(2)}, \dots, W^{(L)}, b^{(L)})$. At initialization, each entry of W^{ℓ} or b^{ℓ} is sampled from i.i.d Gaussians $\mathcal{N}(0, 1)$. Hence, θ_0 is a random variable.

3 Challenges

Now there are two elements that make the model hard to deal with using traditional tools.

1. **The kernel is stochastic.** For the models studied θ_0 is a random variable. Hence $K(\cdot, \cdot; \theta_0)$ is not deterministic.
2. **The kernel is parametrized.** The kernel depends on a parameter θ that varies which is itself being updated during training. Hence $K(\cdot, \cdot; \theta)$ is a kernel that evolves with the training. This reflects making the Eq. 6 not linear.

Solutions that are proposed in the paper:

1. (Theorem 1) In probability, $K(x, y; \theta_0) \rightarrow K_0(x, y)$, i.e. where K_0 is a deterministic kernel.
2. (Theorem 2) Uniformly on t , $K(x, y; \theta_t) \rightarrow K_0(x, y)$ for all $t \in [0, T]$.

4 Relation with other models

We have presented another linear model approximation before, the one that assumed that $\theta = \theta_0 + \beta$ and the number of parameters is so large that training effectively only changes the parameter by a small amount. Then β is small and:

$$f(z; \theta) \approx f(x; \theta_0) + \nabla_{\theta} f(x; \theta_0)^{\top} \beta.$$

In the case, $f(z; \theta_0)$ this would come down to the map: $x \mapsto \nabla_{\theta} f(x; \theta_0)$ followed by the estimation of β using a linear parameter. Notice that this nonlinear map coincides with the nonlinear map we are considering here.

References

- M. Belkin, D. Hsu, S. Ma, and S. Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854, Aug. 2019. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.1903070116.
- A. Jacot, F. Gabriel, C. Hongler. Neural Tangent Kernel: Convergence and Generalization in Neural Networks In *Advances in Neural Information Processing Systems 32*, 2018.
- A. Rahimi and B. Recht. Random Features for Large-Scale Kernel Machines. In *Advances in Neural Information Processing Systems 20*, pages 1177–1184. 2008.